



UNIVERSITÀ DI PARMA

Features Extraction

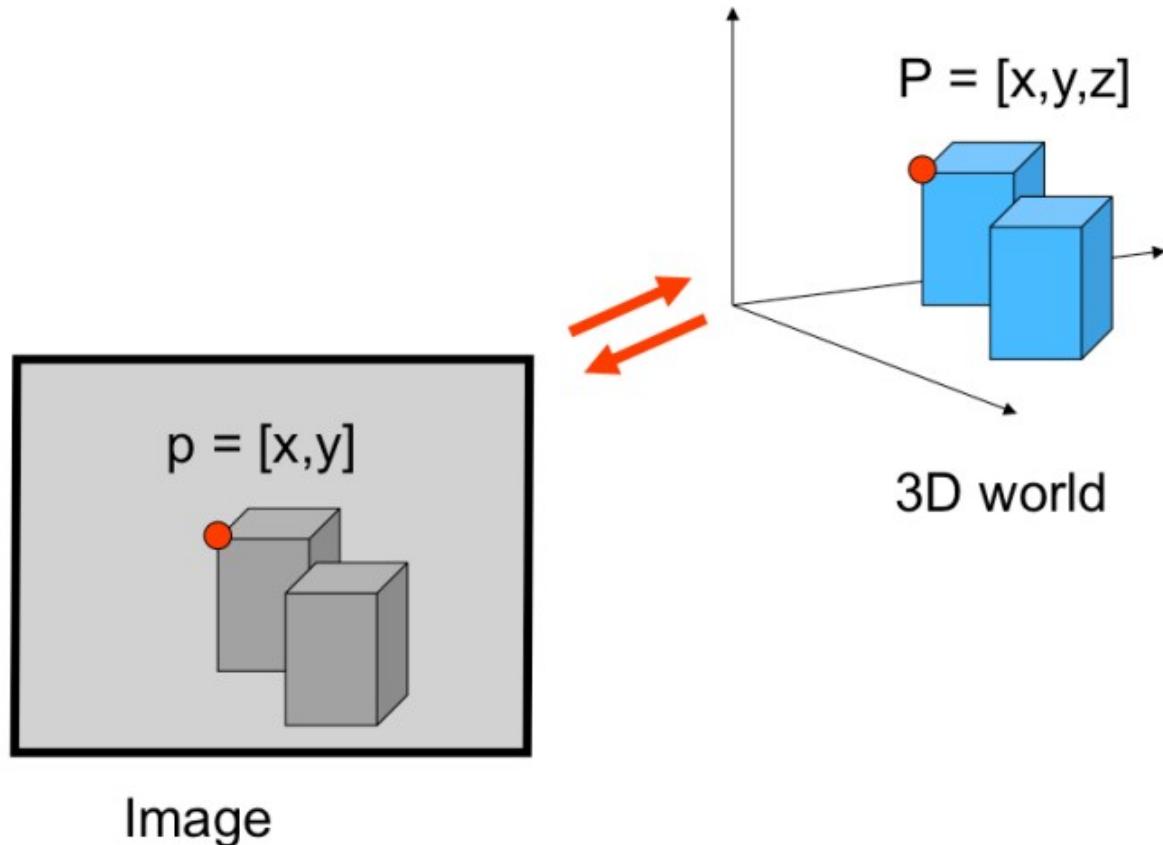


Summary

- What is a feature?
- Keypoints ideal characteristics
- Harris Corners
- Scale Invariant Detectors

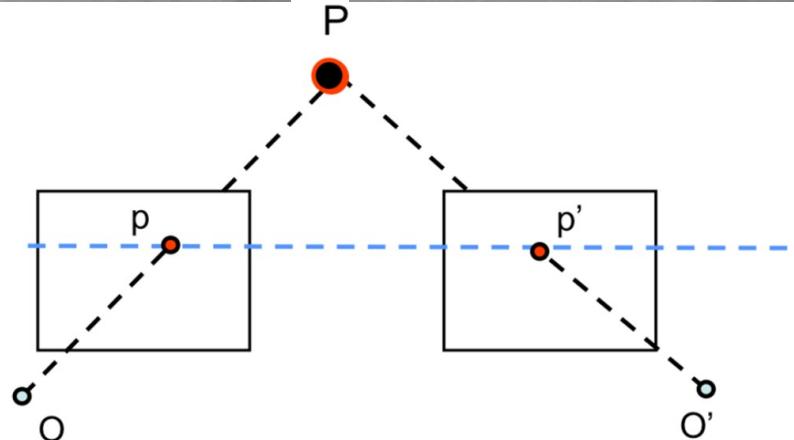
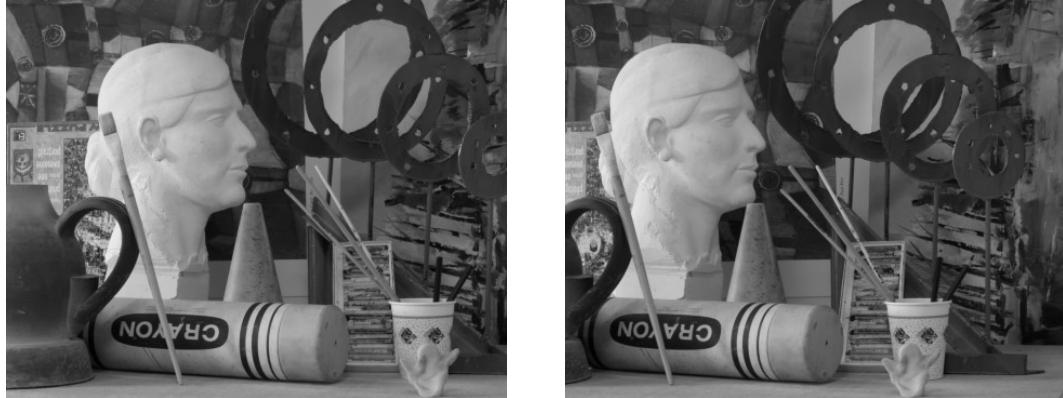
- [FP] D. A. Forsyth and J. Ponce. **Computer Vision: A Modern Approach** (2nd Edition). Prentice Hall, 2011.
- **CS231A · Computer Vision: from 3D reconstruction to recognition**, Prof. Silvio Savarese – Stanford University
- **CS131 · Computer Vision: Foundations and Applications**, Prof. Fei-Fei Li – Stanford University
- **Elementi di Analisi per Visione Artificiale**
 - Paolo Medici <http://www.ce.unipr.it/people/medici>

Small Recap



- We found a geometrical relation between world and image

Small Recap



- We found a geometrical relation between world and image

Beyond Stereo Vision

UNIVERSITÀ
DI PARMA



Beyond Stereo Vision

UNIVERSITÀ
DI PARMA



by [Diva Sian](#)

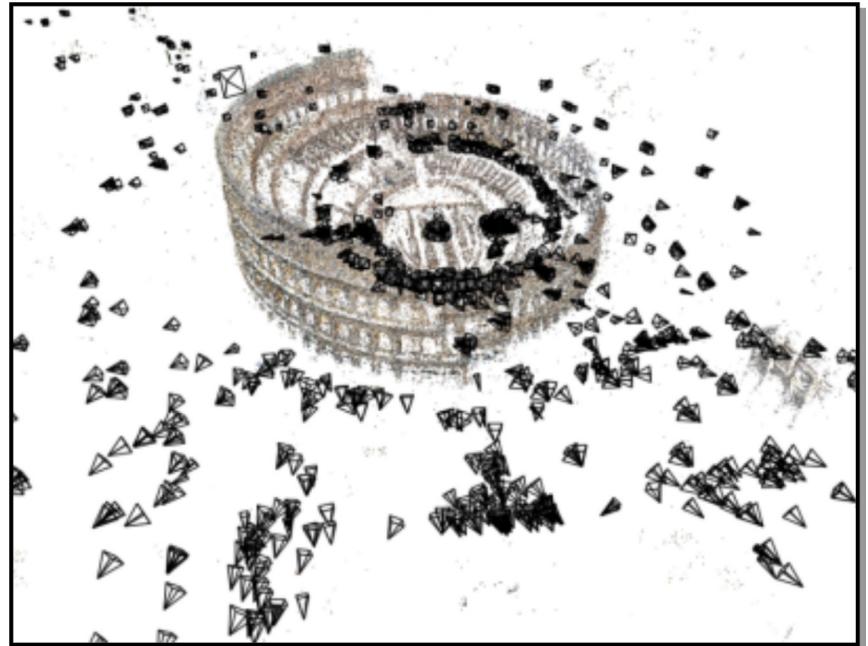
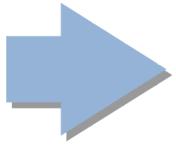


Beyond Stereo Vision

UNIVERSITÀ
DI PARMA

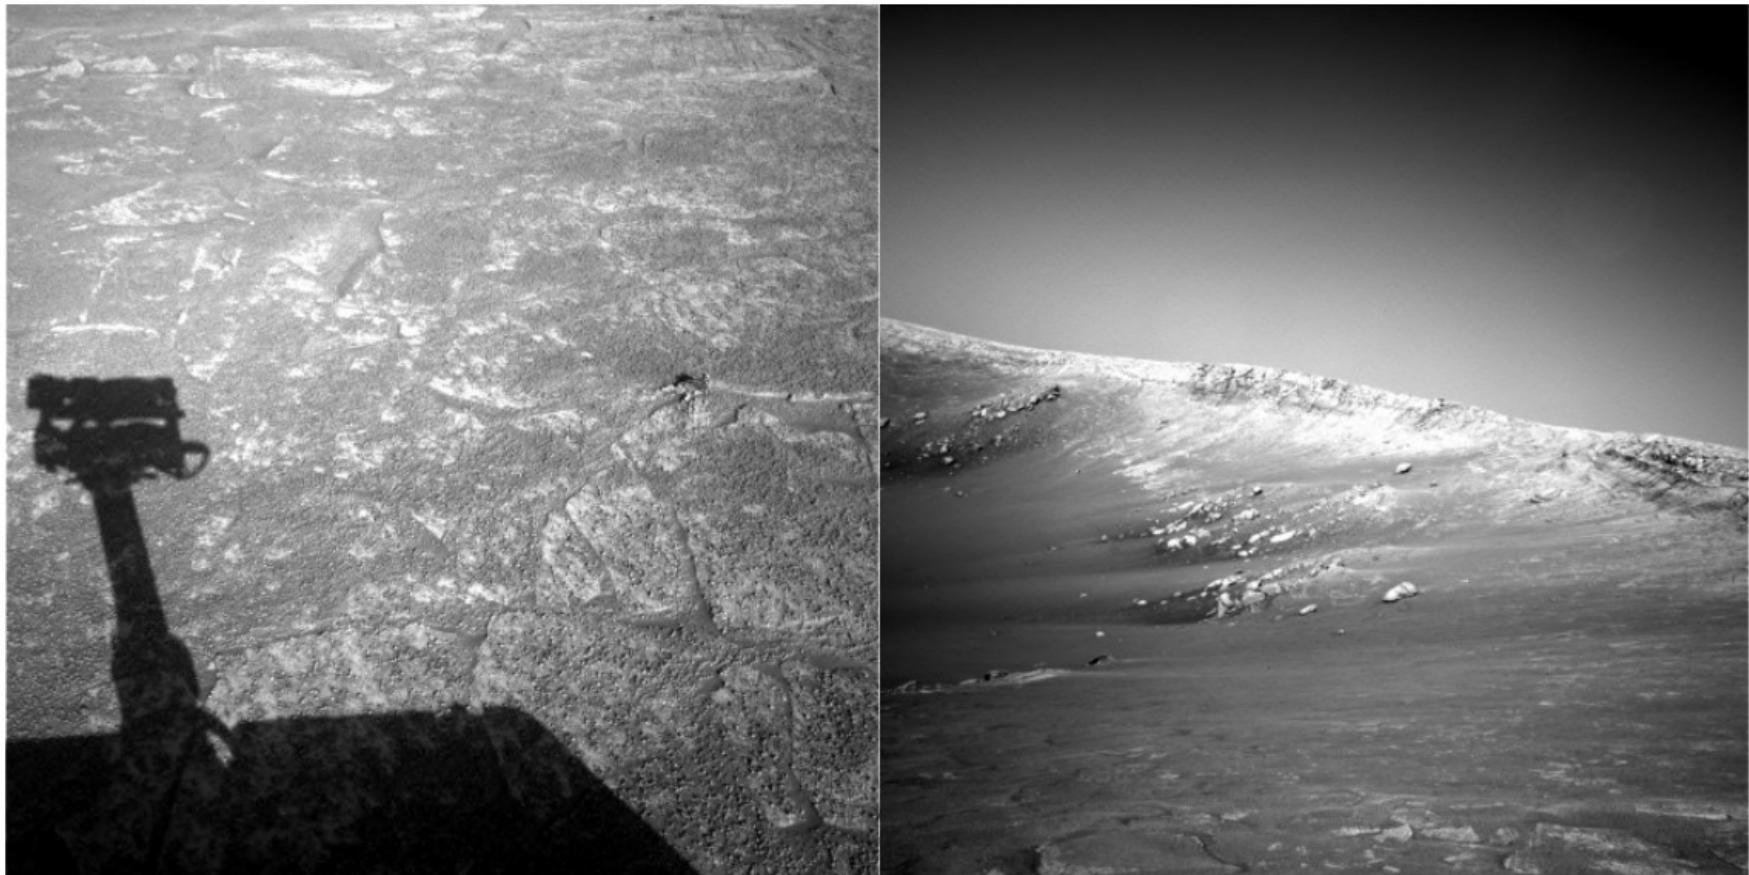


Beyond Stereo Vision



Beyond Stereo Vision

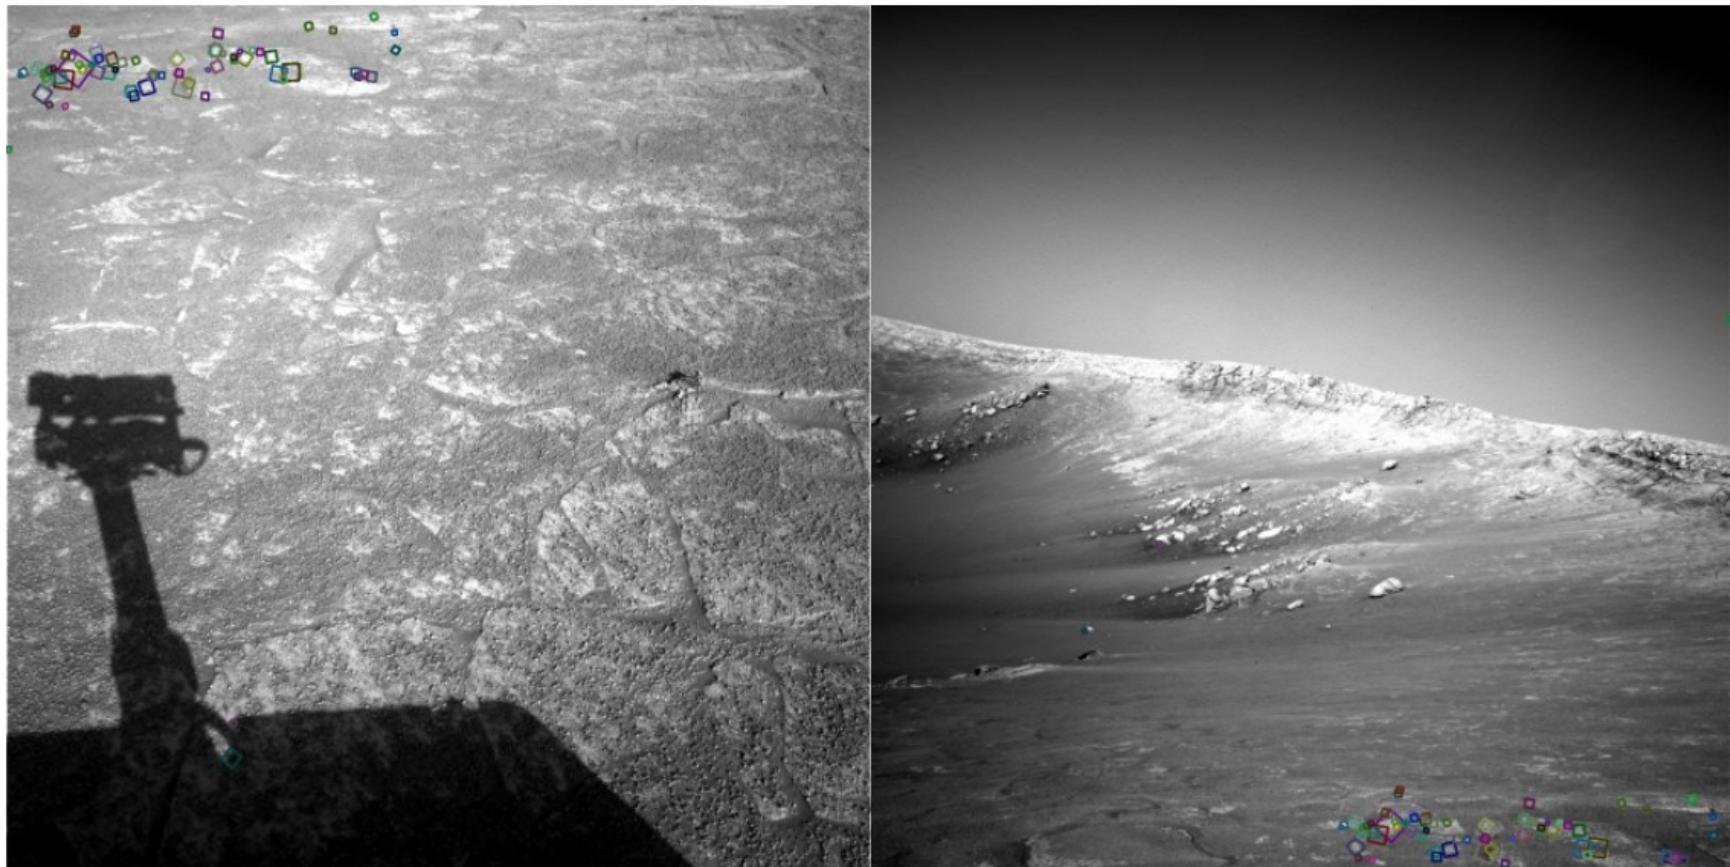
UNIVERSITÀ
DI PARMA



NASA Mars Rover images

Beyond Stereo Vision

UNIVERSITÀ
DI PARMA

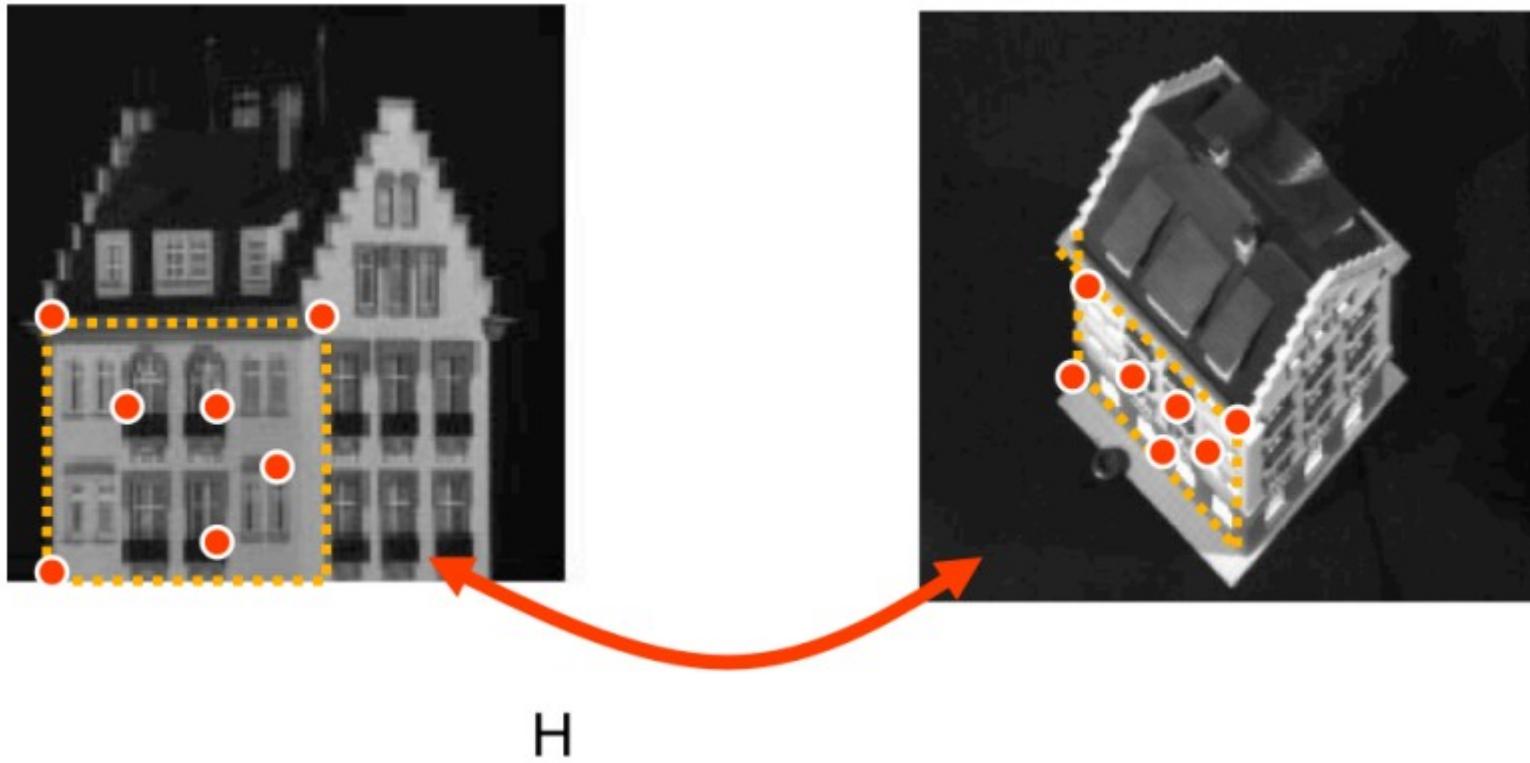


NASA Mars Rover images with SIFT feature matches

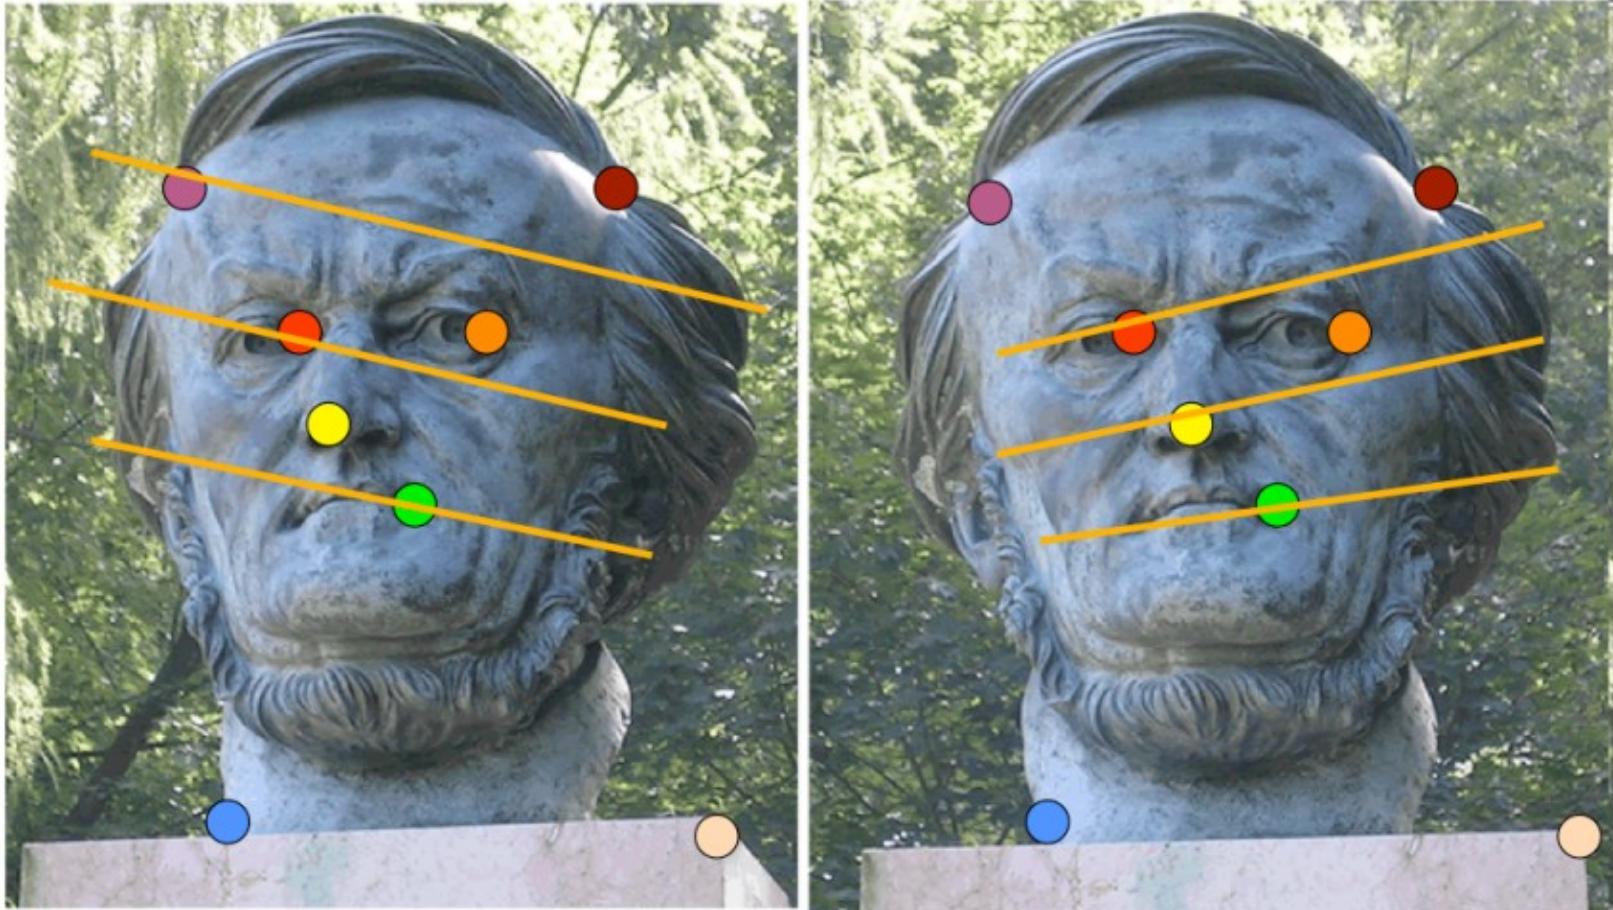
- During this lesson and the following one we are going to learn:
 - What are “features” and their properties
 - How to extract “features” from images
 - How to match them among all images

- Potential outcomes are:
 - Fit or estimate models that describe an image or a portion of it
 - Match or index images
 - Detect objects or actions from images
 - “Combine” different images

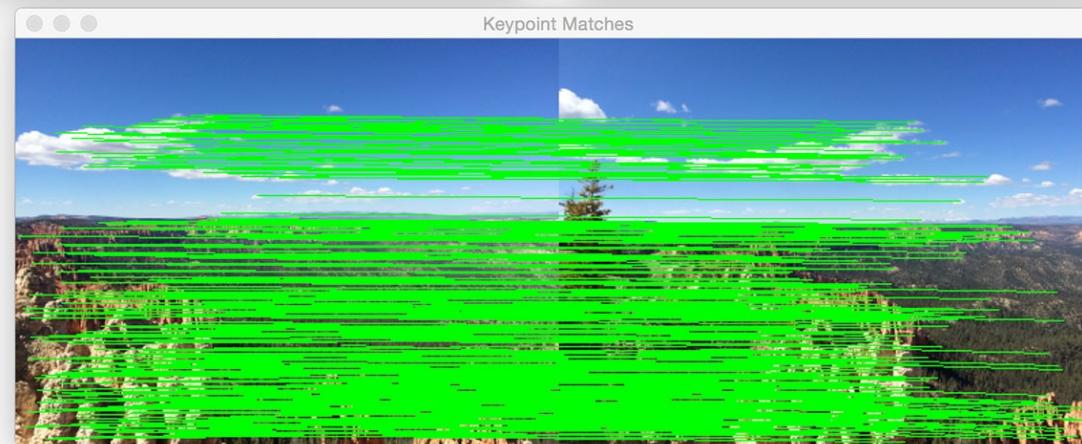
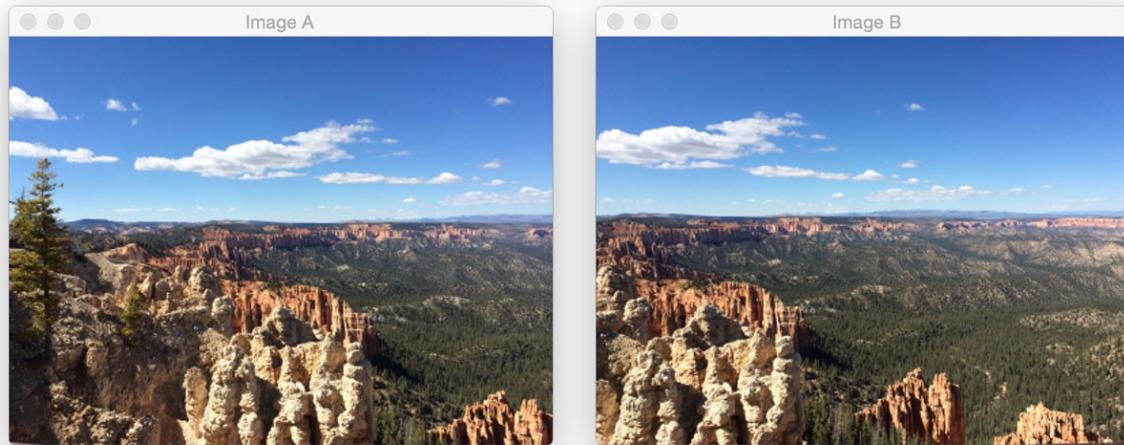
Applications



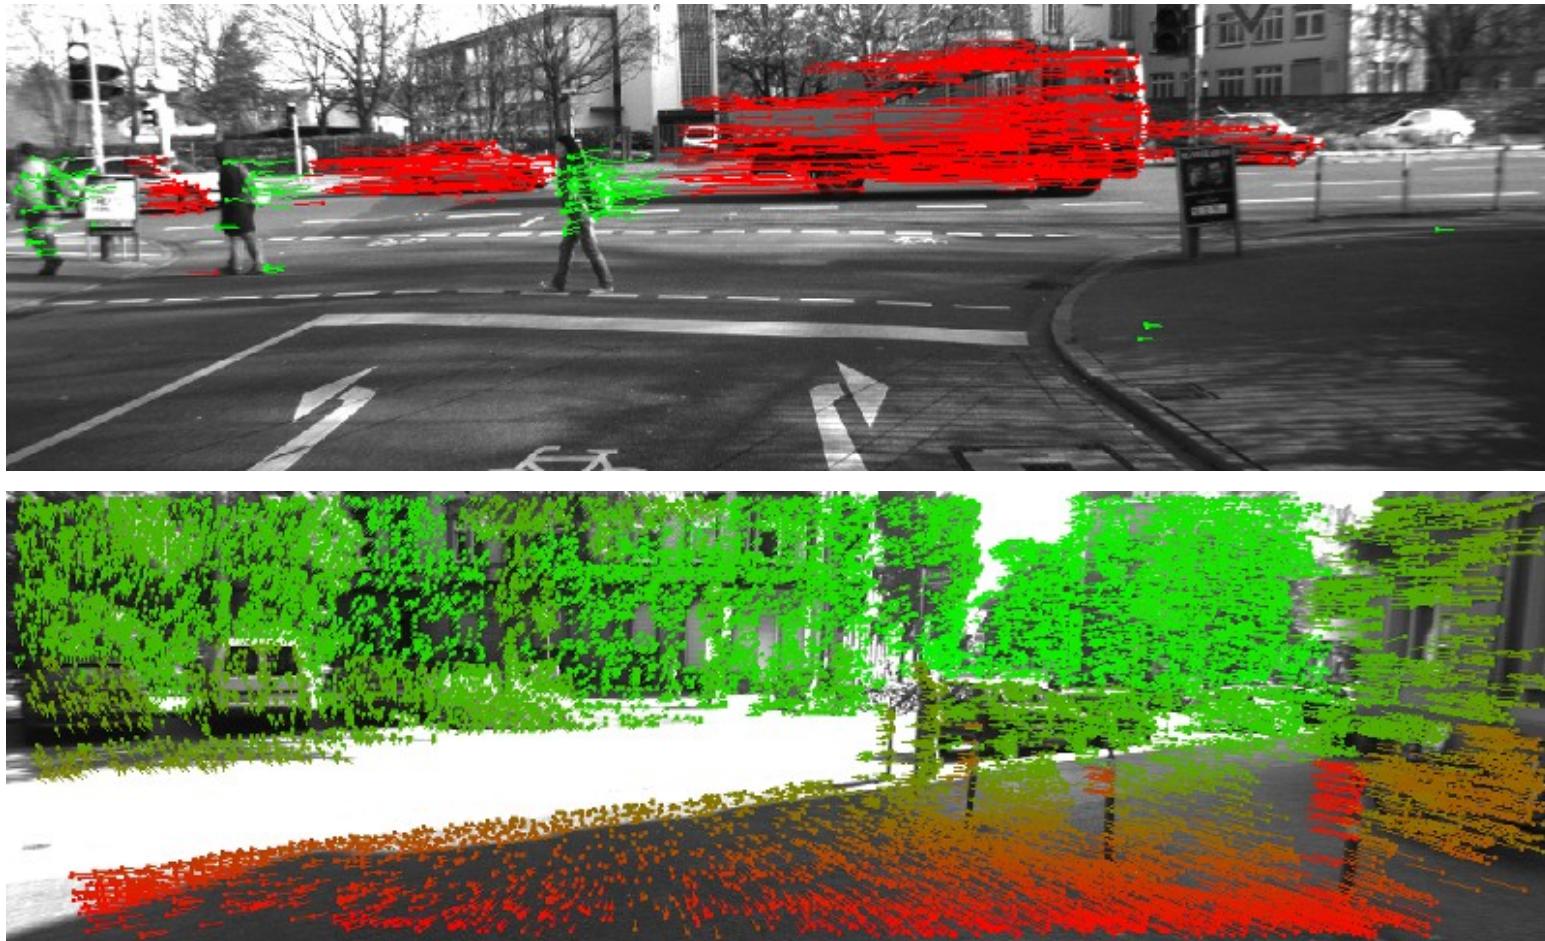
Applications



Applications



Applications



Applications



Applications



How to do this?

- Theoretically we could use stereo vision techniques
 - Too much expensive!
 - Moreover, we do not have to obtain a full 3D reconstruction
- Just look at the following case
 - Joining of different images...

Example



- How to combine following images?



Example



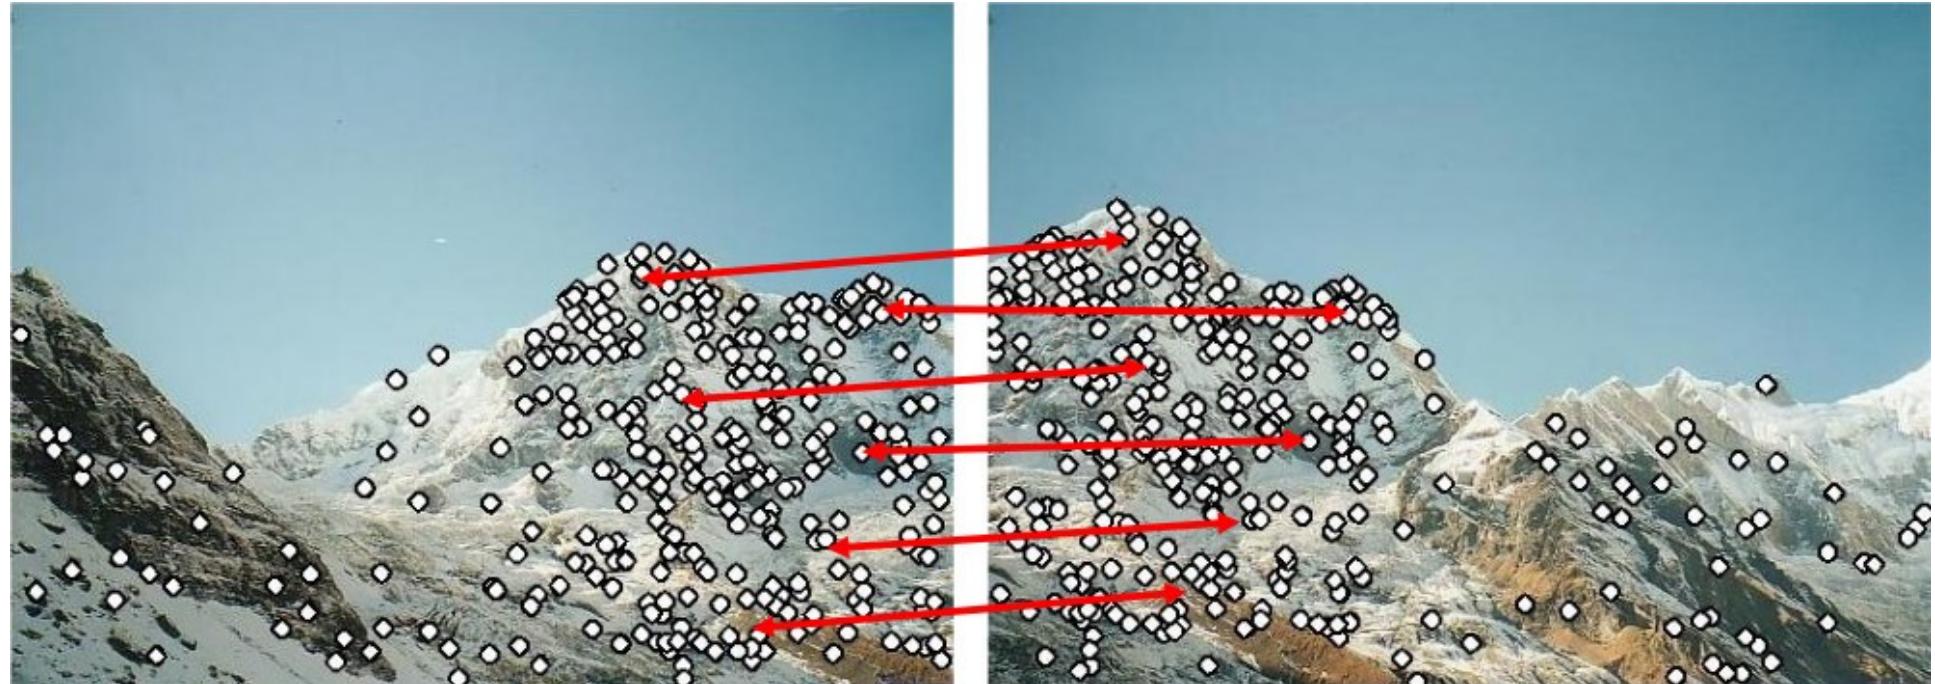
- 1. select “specific” points



Example



- 2. find potential correspondances



Example



- 3. compute alignment



Example recap

- Select specific points → keypoints or features → **features extraction**
- Find potential correspondences → **features matching**
- Application dependent step
 - Matching
 - Indexing
 - Detection
 - Image alignment

Keypoints or Features

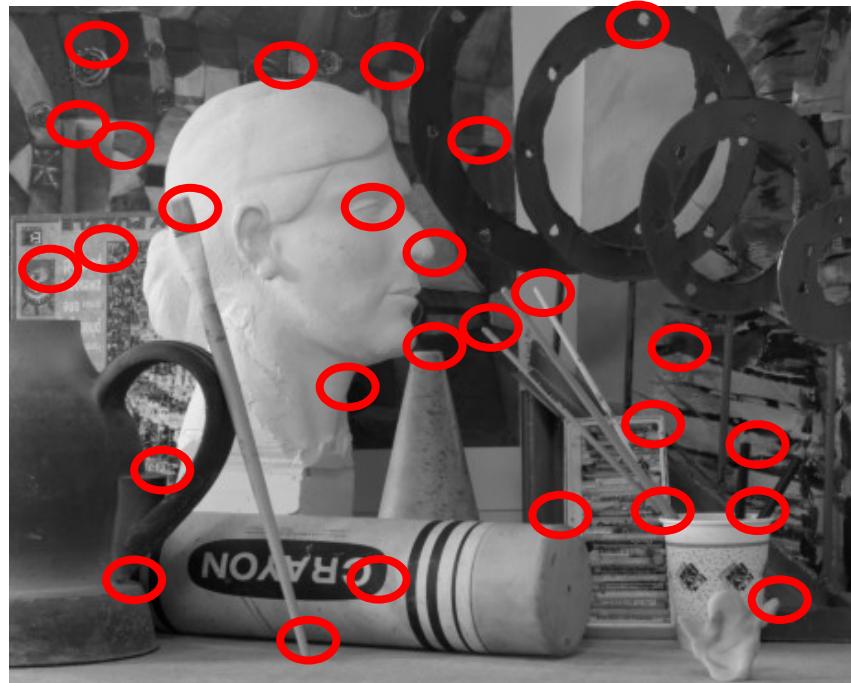


- For Stereo Vision
 - Window match → SAD
- In a more general situation
 - Different size
 - Different orientation
- A complete match is impracticable
 - Only keypoints → patches

Keypoints or Features



- What keypoints I have to select?



Characteristics of a good feature

- **Repeatability:** to be able to find the same features on even very different images
- **Saliency:** the feature contains information
- **Locality:** relatively small area of the image, namely more robust wrt occlusions and clutter

Repeatability



Illumination
invariance



Scale
invariance

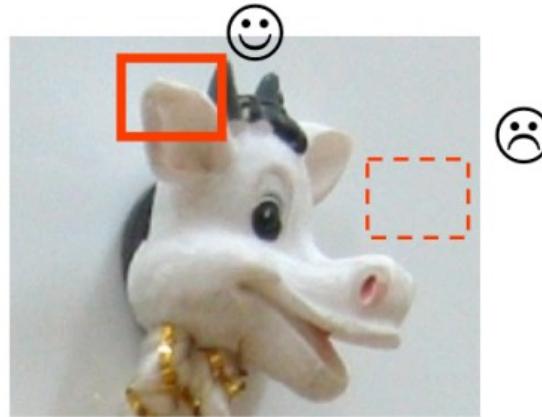


Pose invariance
•Rotation
•Affine

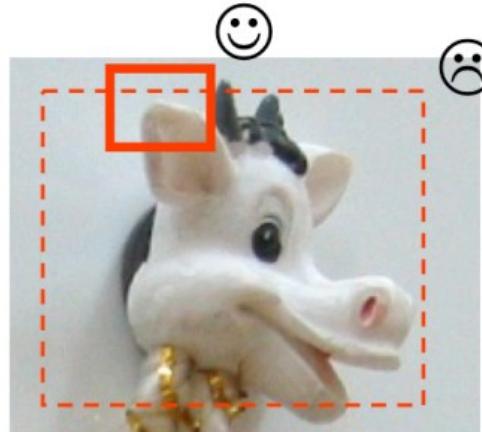
Saliency & Locality



- Saliency

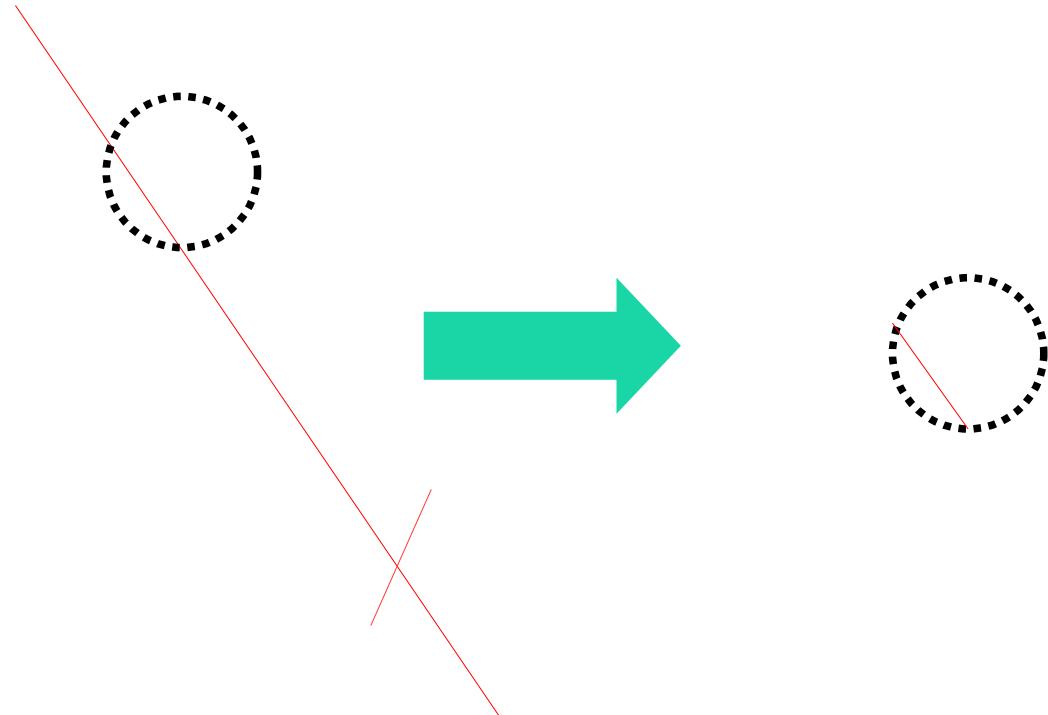


- Locality

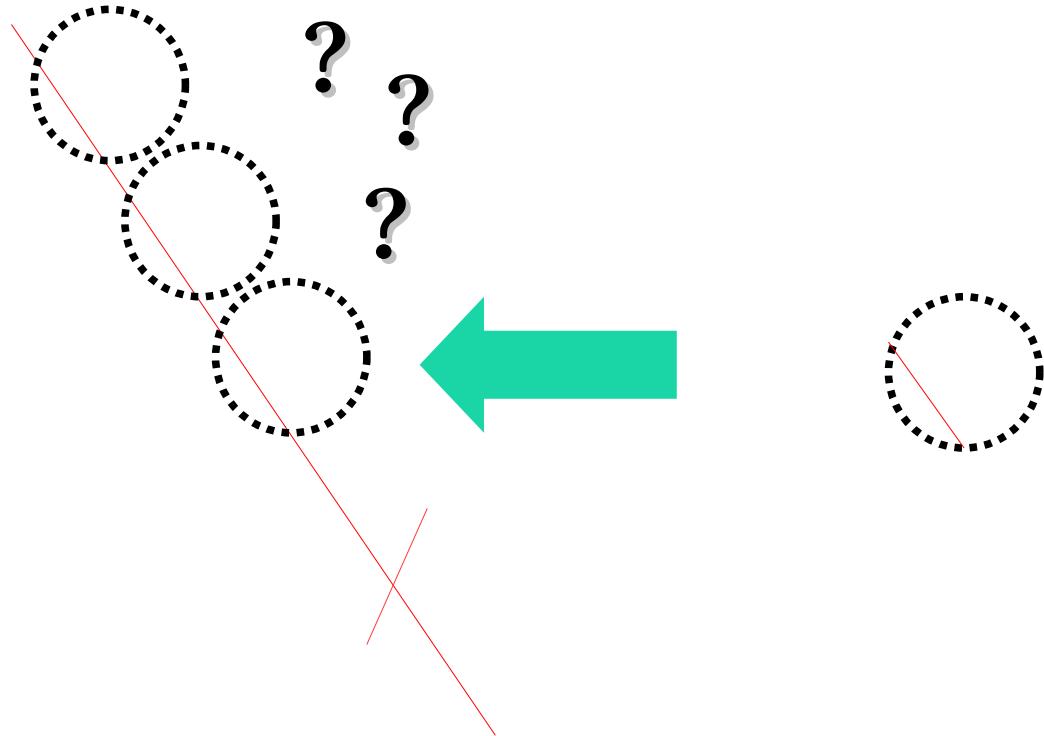


- Textureless patches are nearly impossible to match
- Patches with large contrast changes (gradients) are easier to localize
 - We want edges!
- Are straight edges a good option?
 - Aperture Problem

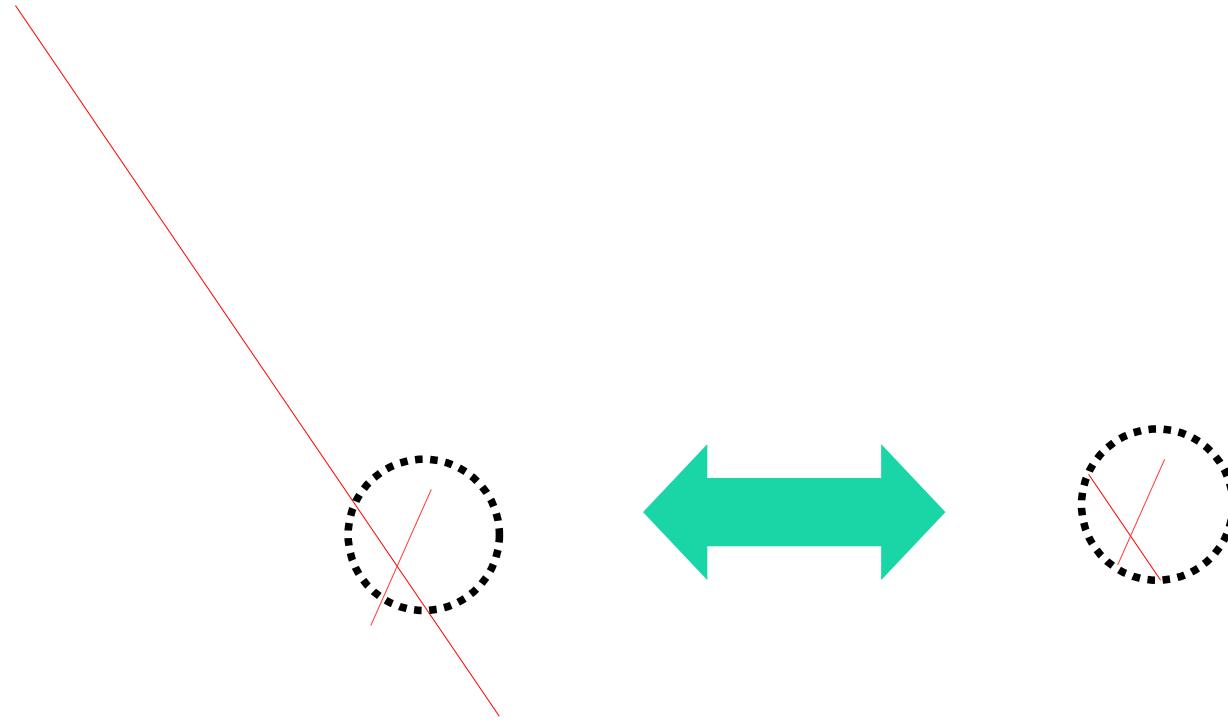
Aperture problem



Aperture problem



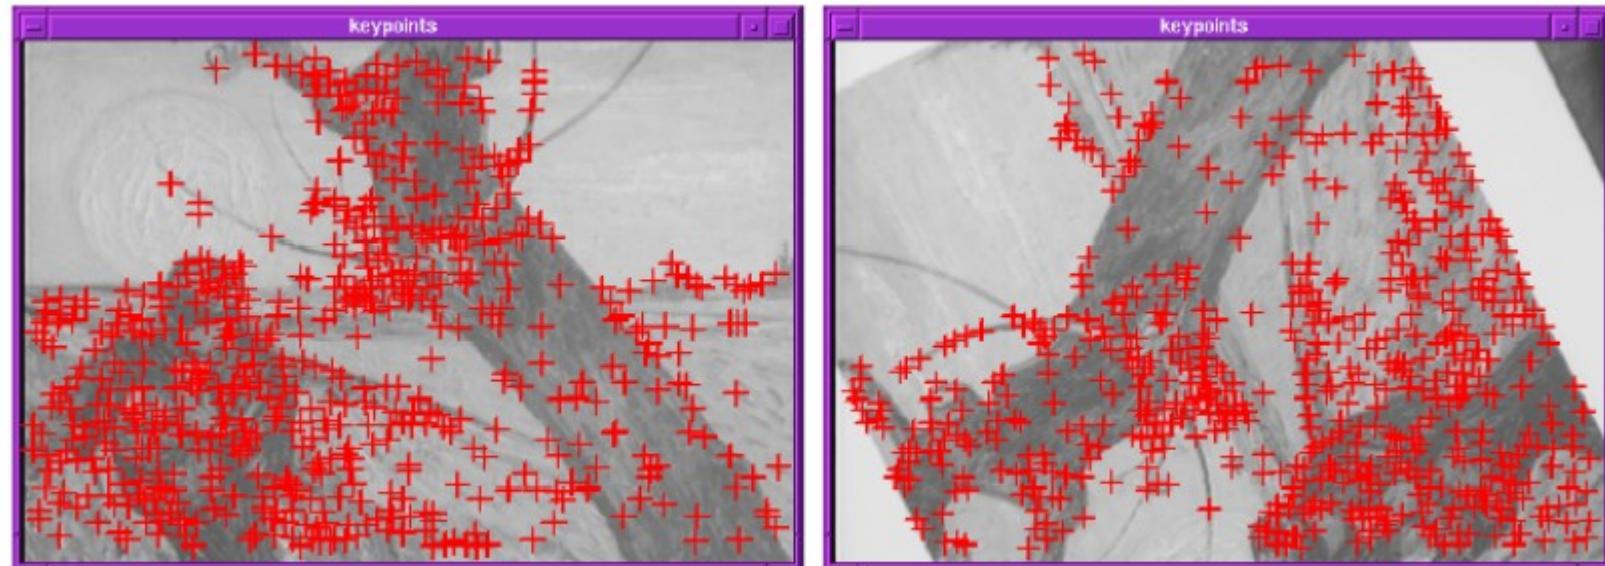
Aperture problem



Corner Detection



- In “real” images a corner is easier to match
 - 2 “strong” gradients

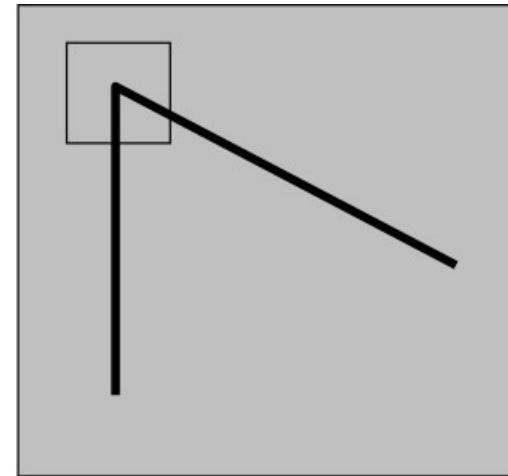
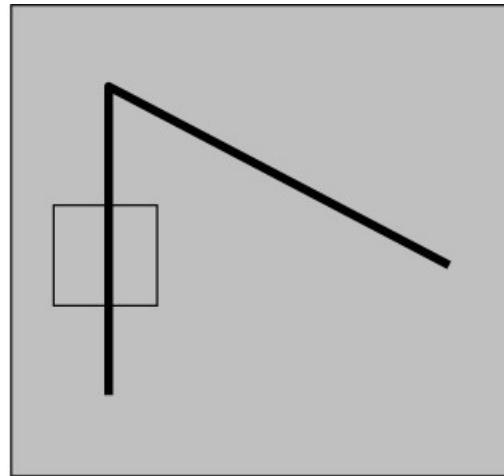
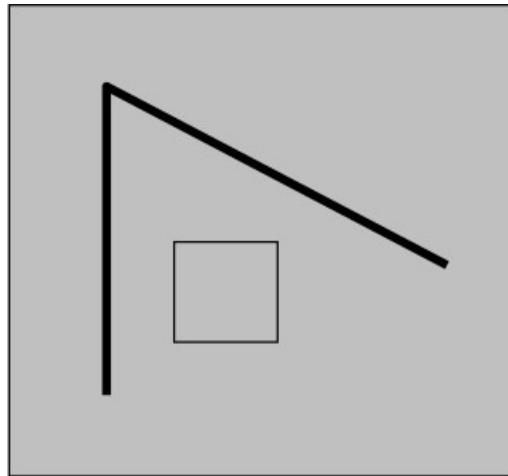


Repeatability – Saliency – Locality

Local Measure of Uniqueness

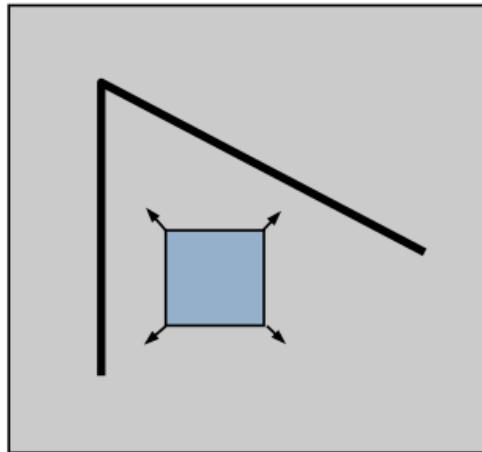


- What defines whether our patch is a good or bad candidate?

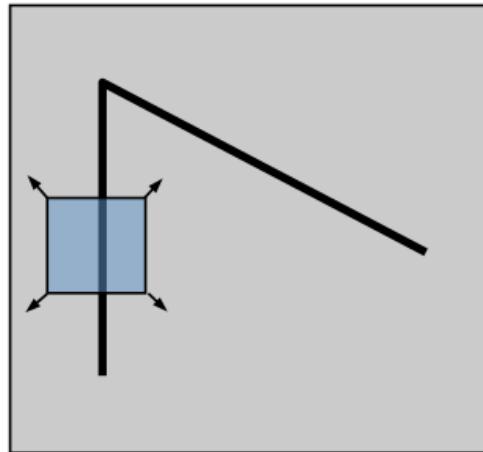


Corner Detection

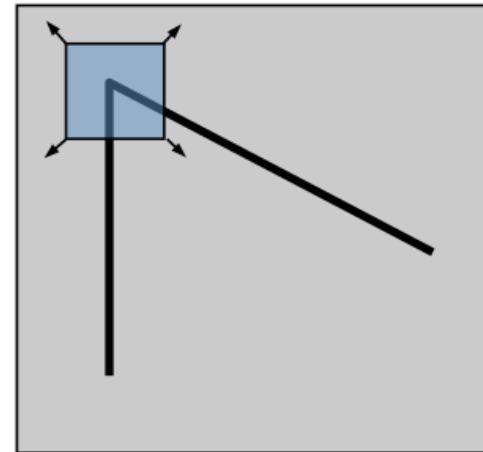
- Let's move a small window (W) on the image
- How the content is changing?



"flat" region:
no change in all
directions



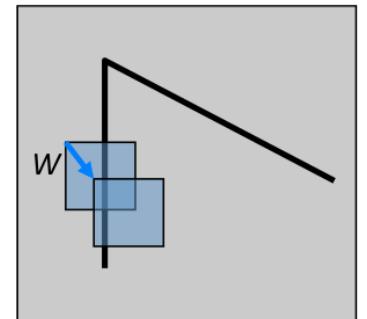
"edge":
no change along the
edge direction



"corner":
significant change in
all directions

- Let's move window W by (u,v)
 - u and v are small...
- A good estimation can be the Sum of Squared Differences (SSD)
- The SSD error $E(u, v)$ can be defined as

$$E(u, v) = \sum_{(x, y) \in W} [I(x+u, y+v) - I(x, y)]^2$$



Corner Detector

- Using a Taylor Expansion

$$I_x = \frac{\partial I(x, y)}{\partial x} \quad I_y = \frac{\partial I(x, y)}{\partial y}$$

$$I(x+u, y+v) = I(x, y) + \frac{I_x u}{1!} + \frac{I_y v}{1!} + \text{higher order terms}$$

$$\approx I(x, y) + I_x u + I_y v$$

$$\approx I(x, y) + [I_x \quad I_y] \begin{bmatrix} u \\ v \end{bmatrix}$$

- Using a Taylor Expansion

$$I_x = \frac{\partial I(x, y)}{\partial x} \quad I_y = \frac{\partial I(x, y)}{\partial y}$$

$$E(u, v) = \sum_{(x, y) \in W} [I(x+u, y+v) - I(x, y)]^2$$

$$\approx \sum_{(x, y) \in W} [I(x, y) + I_x u + I_y v - I(x, y)]^2$$

$$\approx \sum_{(x, y) \in W} [I_x u + I_y v]^2$$

Corner Detector

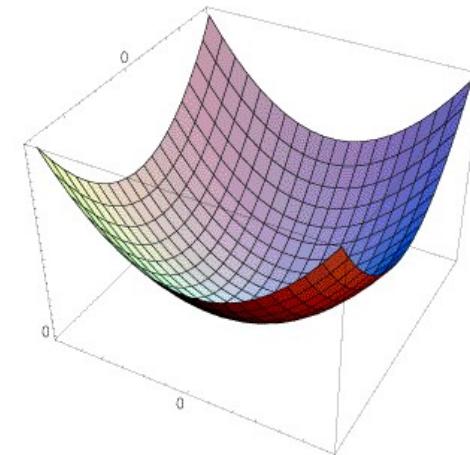


- Using a Taylor Expansion

$$I_x = \frac{\partial I(x, y)}{\partial x} \quad I_y = \frac{\partial I(x, y)}{\partial y}$$

$$E(u, v) \approx \sum_{(x, y) \in W} [I_x u + I_y v]^2$$

$$\approx A u^2 + 2 B u v + C v^2$$



$$A = \sum_{(x, y) \in W} I_x^2$$

$$B = \sum_{(x, y) \in W} I_x I_y$$

$$C = \sum_{(x, y) \in W} I_y^2$$

Corner Detector

- Using a Taylor Expansion

$$I_x = \frac{\partial I(x, y)}{\partial x} \quad I_y = \frac{\partial I(x, y)}{\partial y}$$

$$E(u, v) \approx \sum_{(x, y) \in W} [I_x u + I_y v]^2$$

$$\approx A u^2 + 2 B u v + C v^2$$

$$\approx \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} u & v \end{bmatrix} H \begin{bmatrix} u \\ v \end{bmatrix}$$

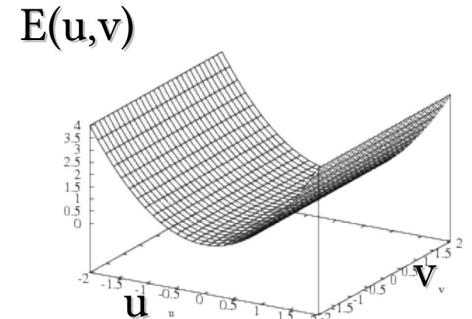
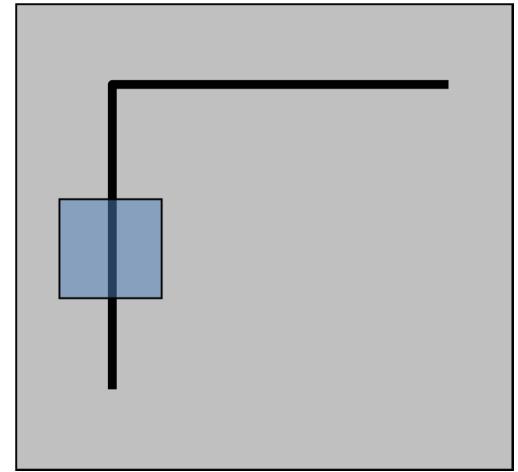
Corner Detector



- Vertical Edge $\rightarrow I_y=0$

$$E(u, v) \approx [u \quad v] \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$\approx [u \quad v] \begin{bmatrix} A & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

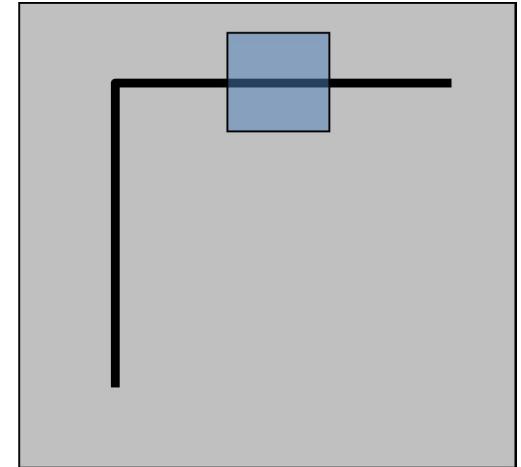


Corner Detector

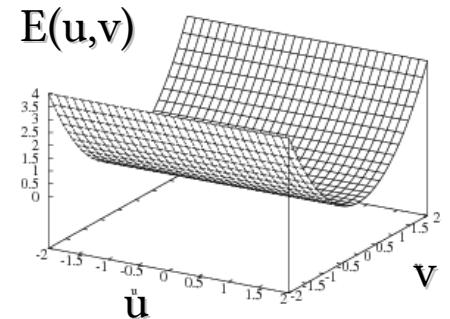


- Horizontal Edge $\rightarrow I_x=0$

$$E(u, v) \approx [u \quad v] \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



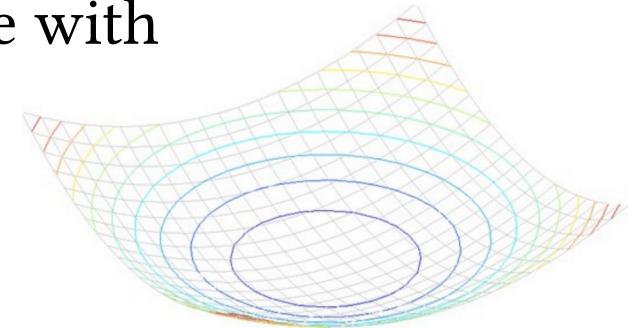
$$\approx [u \quad v] \begin{bmatrix} 0 & 0 \\ 0 & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$



- General Case

$$E(u, v) \approx [u \quad v] \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = [u \quad v] H \begin{bmatrix} u \\ v \end{bmatrix}$$

- H can be considered as describing an ellipse with
 - Axis length \propto H eigenvalues
 - Axis orientation \rightarrow H eigenvectors



Eigenvector/Eigenvalues review

- Given a matrix A eigenvectors x and eigenvalues λ satisfy:

$$Ax = \lambda x$$

- Eigenvalues can be find solving

$$\det(A - \lambda I) = 0$$

- For a 2×2 matrix like H we can find $\lambda \rightarrow$
 - This leads to a **quadratic** equation
 - i.e. λ_1 and λ_2
- Once found λ s, x can be computed as \rightarrow

$$\det \begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{12} & h_{22} - \lambda \end{bmatrix}$$

$$\begin{bmatrix} h_{11} - \lambda & h_{12} \\ h_{12} & h_{22} - \lambda \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$

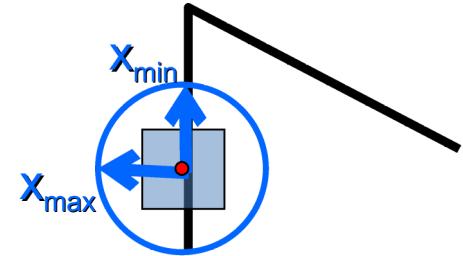
Corner Detector



$$E(u, v) \approx [u \quad v] \begin{bmatrix} A & B \\ B & C \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = [u \quad v] H \begin{bmatrix} u \\ v \end{bmatrix}$$

$$\begin{aligned} Hx_{max} &= \lambda_{max} x_{max} \\ Hx_{min} &= \lambda_{min} x_{min} \end{aligned}$$

- $x_{max} \rightarrow$ direction of largest increase for $E(u, v)$
- $x_{min} \rightarrow$ direction of smallest increase for $E(u, v)$
- $\lambda_{max} \rightarrow$ amount of increase along x_{max}
- $\lambda_{min} \rightarrow$ amount of increase along x_{min}



Corner Detector

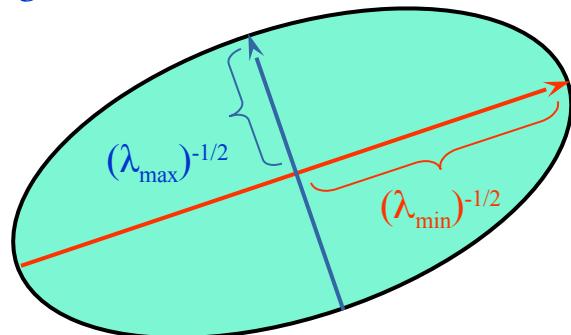


- H can be considered as describing an ellipse with
 - Axis length \propto H eigenvalues
 - Axis orientation \rightarrow H eigenvectors

direction of the fastest
change

$\lambda_{\max}, \lambda_{\min}$: eigenvalues of H

direction of the slowest change

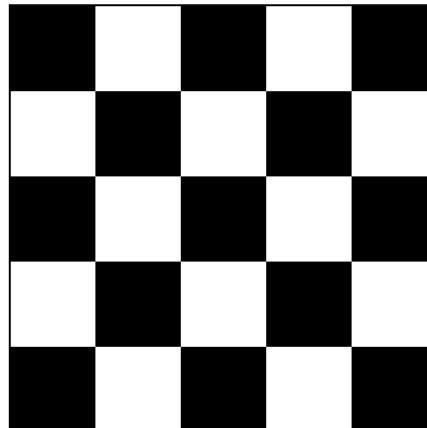


$$H = R \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R^{-1}$$

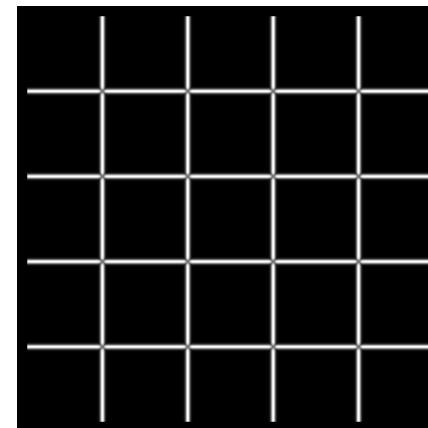
Corner Detector



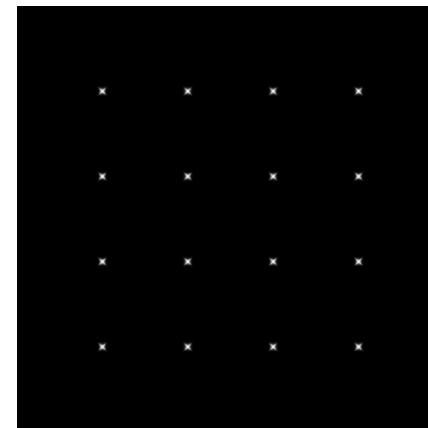
- How λ_{\min} and λ_{\max} behave on corners?
- λ_{\max} is “big” when we have at least one “strong” gradient
- λ_{\min} encodes the strength of the orthogonal gradient



I



λ_{\max}



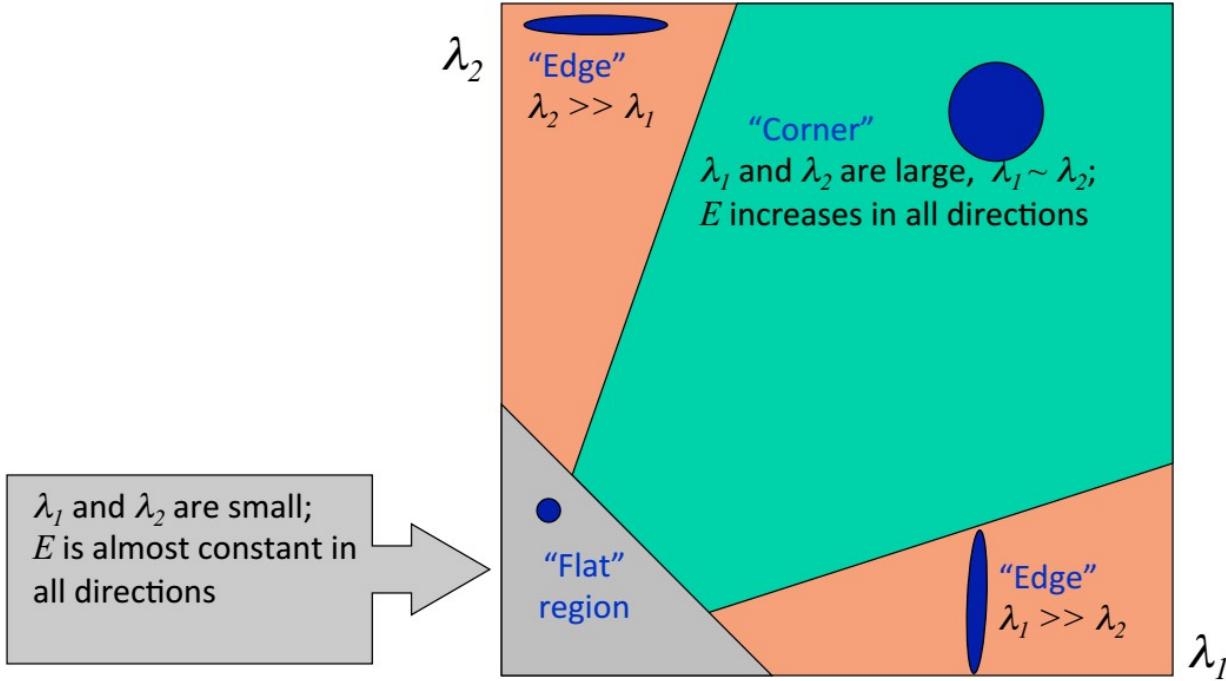
λ_{\min}

Corner Detector



- In a corner both λ_1 and λ_2 are “large”

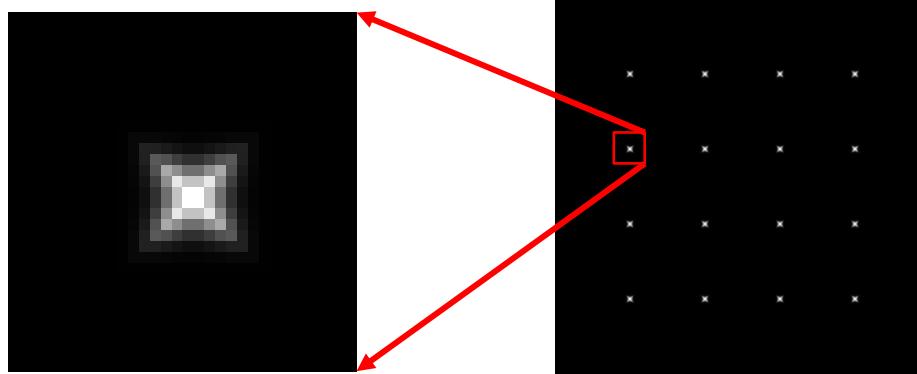
Slide credit: Kristen Grauman



Corner Detector

- Summary
 - Compute horizontal and vertical gradients
 - For each pixel
 - Create H matrix from gradient
 - Compute H eigenvalues
 - Discard pixels with $\lambda_{\min} \ll \lambda_{\max} < \text{threshold}$
 - Apply a NMS approach to find pixels where λ_{\min} is a local maximum

$$H = \begin{bmatrix} \sum_{(x,y) \in W} I_x^2 & \sum_{(x,y) \in W} I_x I_y \\ \sum_{(x,y) \in W} I_x I_y & \sum_{(x,y) \in W} I_y^2 \end{bmatrix}$$

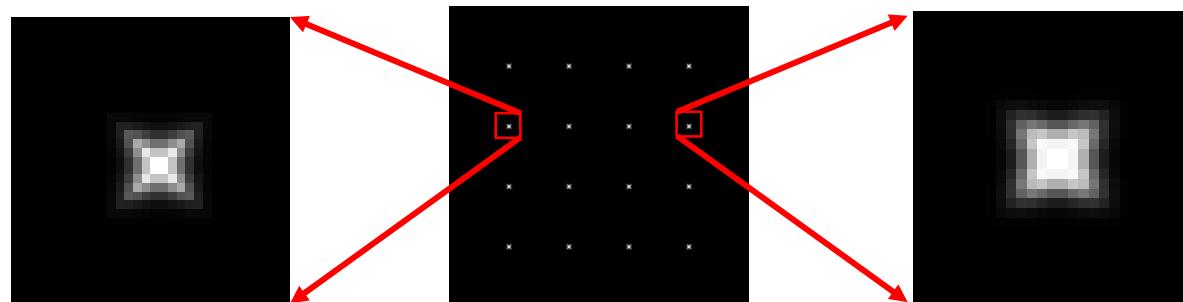


Harris Corner Detector



- “Harris” is a little variant of previous detector
 - No need to compute λ s \rightarrow no square roots
 - Weight function \rightarrow based on the distance from the center pixel

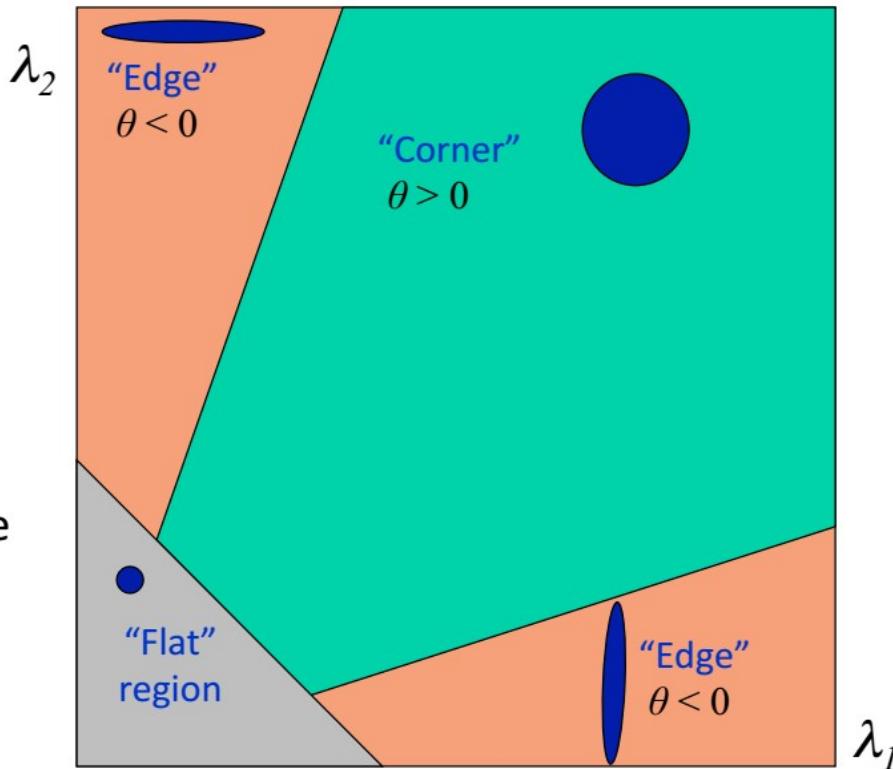
$$\theta = \det(H) - \alpha \operatorname{trace}(H)^2$$



Harris Corner Detector



$$\theta = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$



Slide credit: Kristen Grauman

- Fast approximation
 - Avoid computing the eigenvalues
 - α : constant (0.04 to 0.06)

Corner Detectors

- Shi et al. (1994) → $\lambda_{min}^{\frac{1}{2}}$
- Harris et al. (1998) → $\theta = \det(H) - \alpha \text{trace}(H)^2 = \lambda_0 \lambda_1 - \alpha (\lambda_0 + \lambda_1)^2$
- Triggs et al. (2004) → $\lambda_{min} - \alpha \lambda_{max}$
- Brown et al. (2005) → $\lambda_{min} \approx \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2} = \frac{\det(H)}{\text{trace}(H)}$



- Harris also introduced a window function
 - The idea is to weight distance

$$E(u, v) = \sum_{(x, y) \in W} w(x, y) [I(x+u, y+v) - I(x, y)]^2$$

Window
Function

Shifted
Intensity

Intensity

Window function



$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Option 1: uniform window

- Sum over square window

$$M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

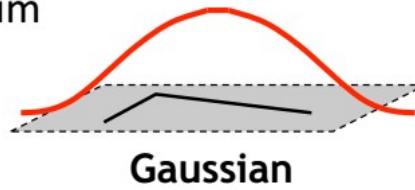


- Problem: not rotation invariant

- Option 2: Smooth with Gaussian

- Gaussian already performs weighted sum

$$M = g(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$



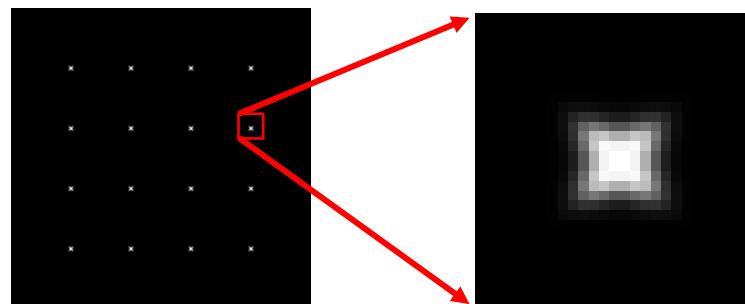
- Result is rotation invariant

Harris Corner Detector



- Summary
 - Compute horizontal and vertical gradients
 - Apply window function
 - For each pixel
 - Create H matrix from gradients
 - Compute θ
 - Discard pixels with $\theta < 0$ or anyway below a given threshold
 - Apply a NMS approach to find pixels where λ_{\min} is a local maximum

$$H(\sigma_I, \sigma_D) = g(\sigma_I) \begin{bmatrix} \sum_{(x,y) \in W} I_x^2(\sigma_D) & \sum_{(x,y) \in W} I_x I_y(\sigma_D) \\ \sum_{(x,y) \in W} I_x I_y(\sigma_D) & \sum_{(x,y) \in W} I_y^2(\sigma_D) \end{bmatrix}$$

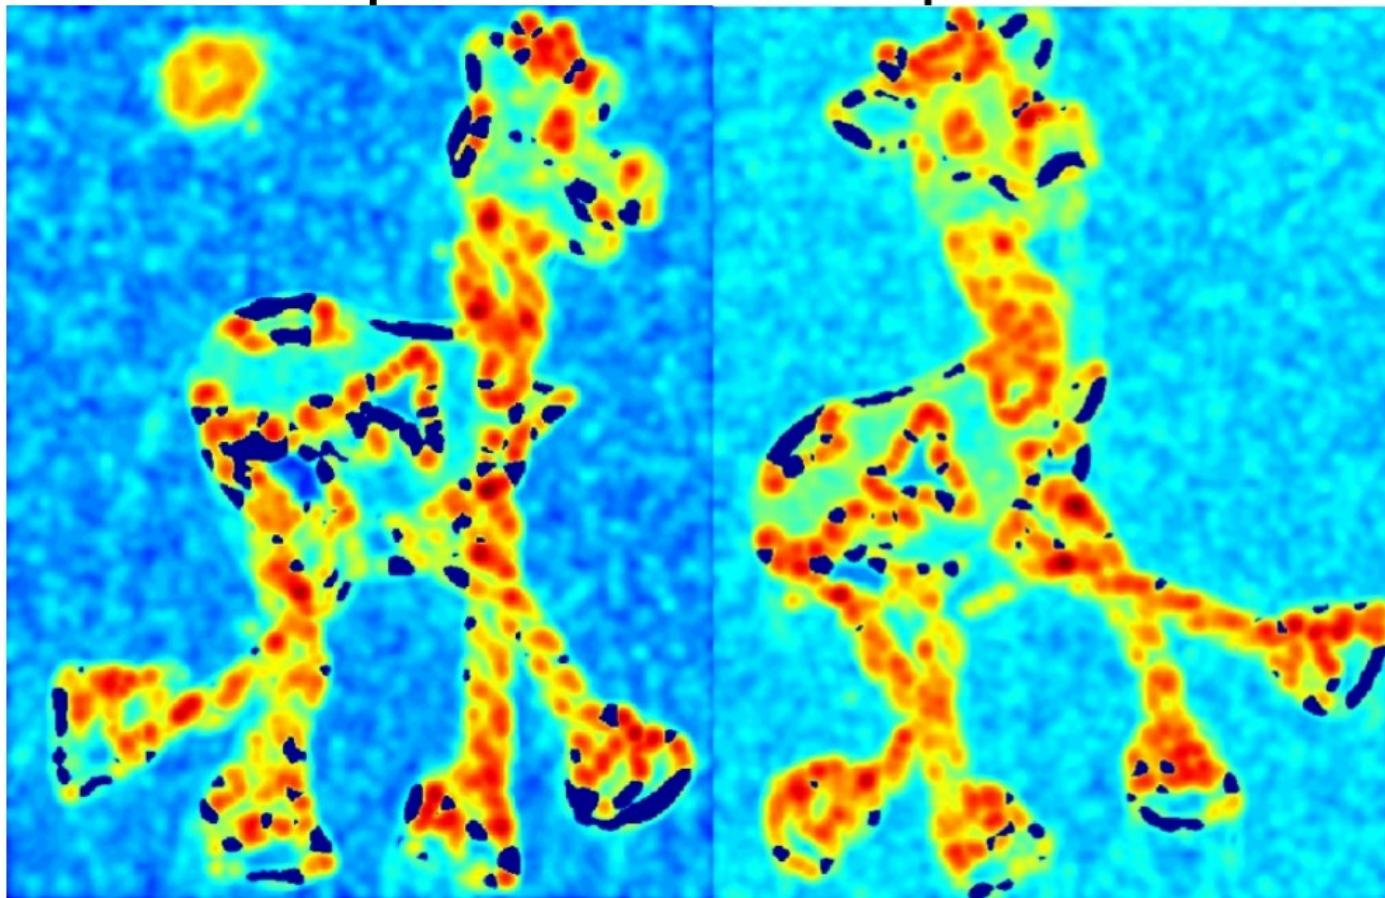


Harris Corner Detector steps



Slide adapted from Darya Frolova, Denis Simakov

Harris Corner Detector steps



Slide adapted from Darya Frolova, Denis Simakov

Harris Corner Detector steps



Harris Corner Detector steps



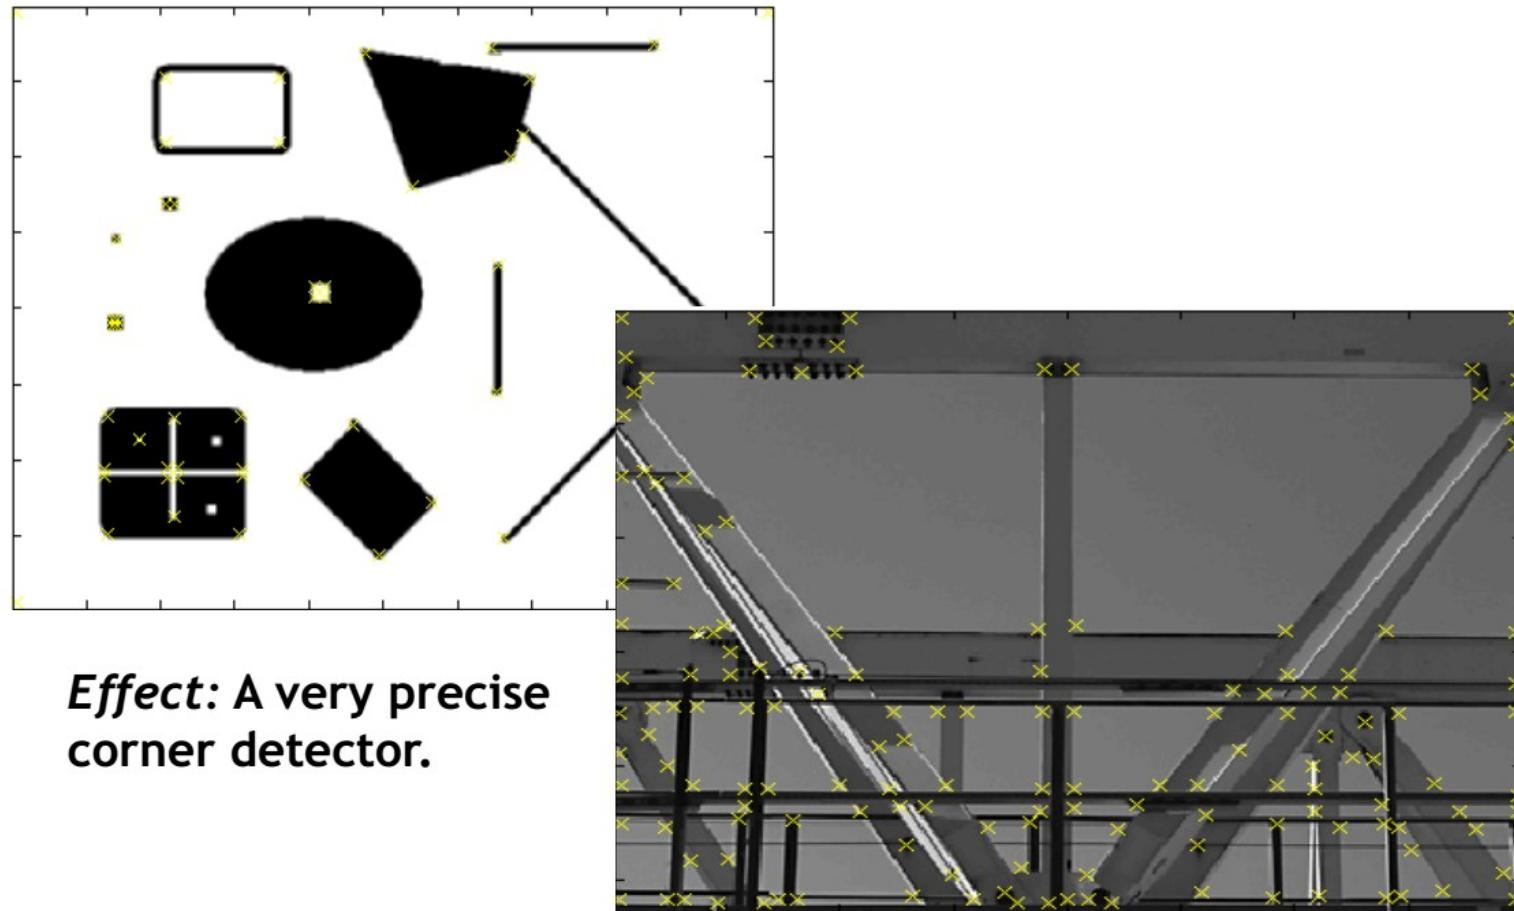
Slide adapted from Darya Frolova, Denis Simakov

Harris Corner Detector steps



Slide adapted from Darya Frolova, Denis Simakov

Examples



Slide credit: Krystian Mikolajczyk

Examples



Slide credit: Krystian Mikolajczyk

Sparse Stereo

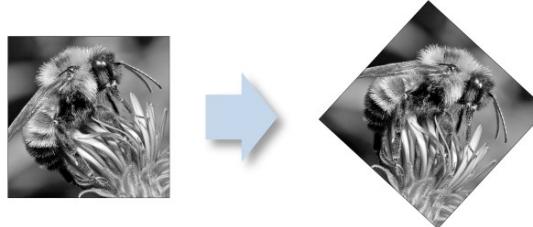


Invariance to transformations

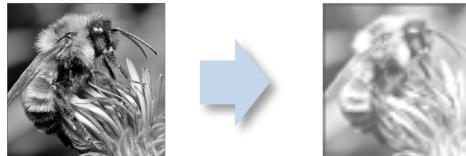


- How much keypoints extraction is robust to image transformations?

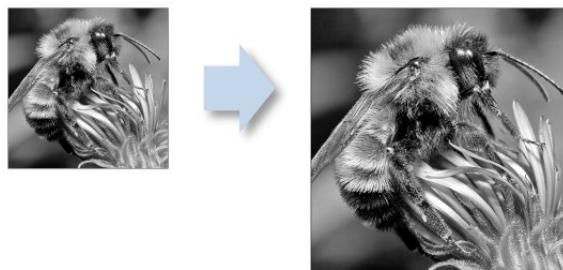
- Rotation (traslation)



- Intensity change

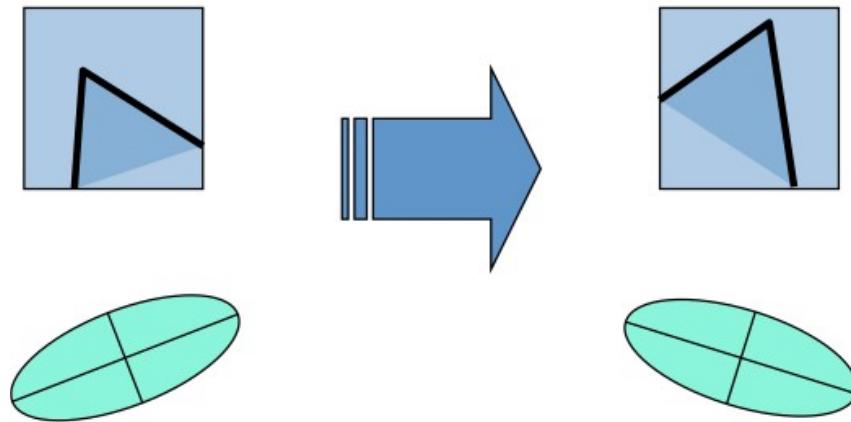


- Scale



- Ellipse rotates but eigenvalues are the same

$$H = R \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R^{-1}$$

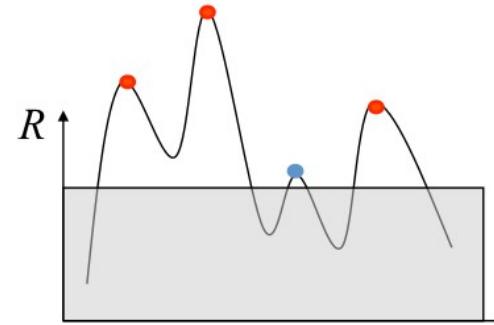
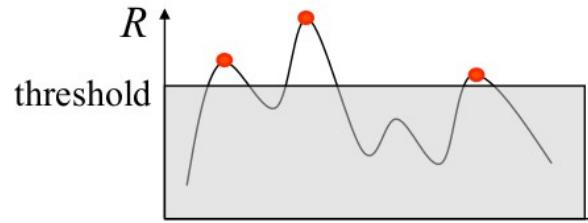


Invariant to Rotation

Intensity Change



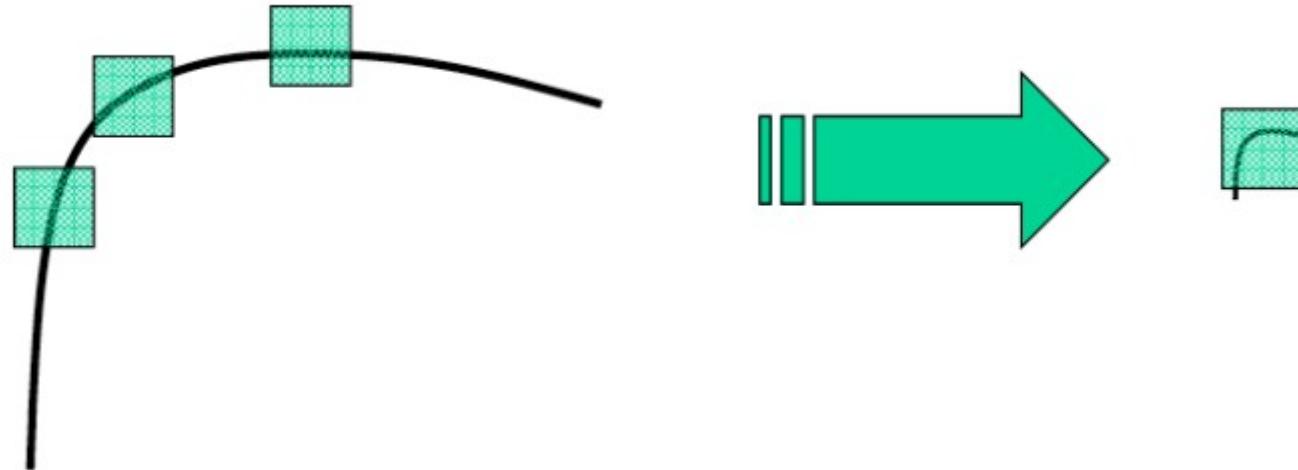
- Intensity change can be modeled as $g(r, c) = a \cdot f(r, c) + b$
- Derivative operations make result invariant wrt b
- Not so true about a ...



Partially Invariant to Intensity Changes



- Right → corner
- Left → corner

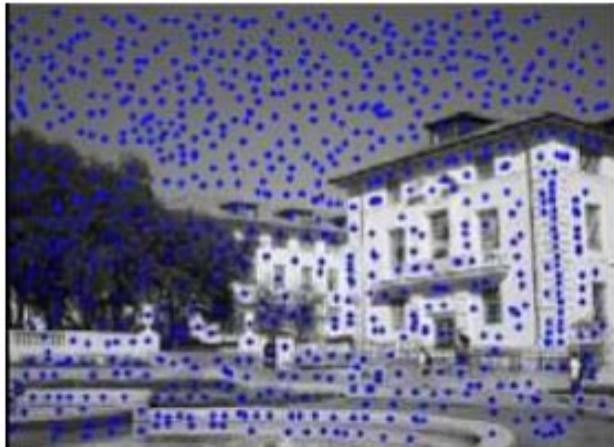


Not Invariant to Scale Changes

Scale



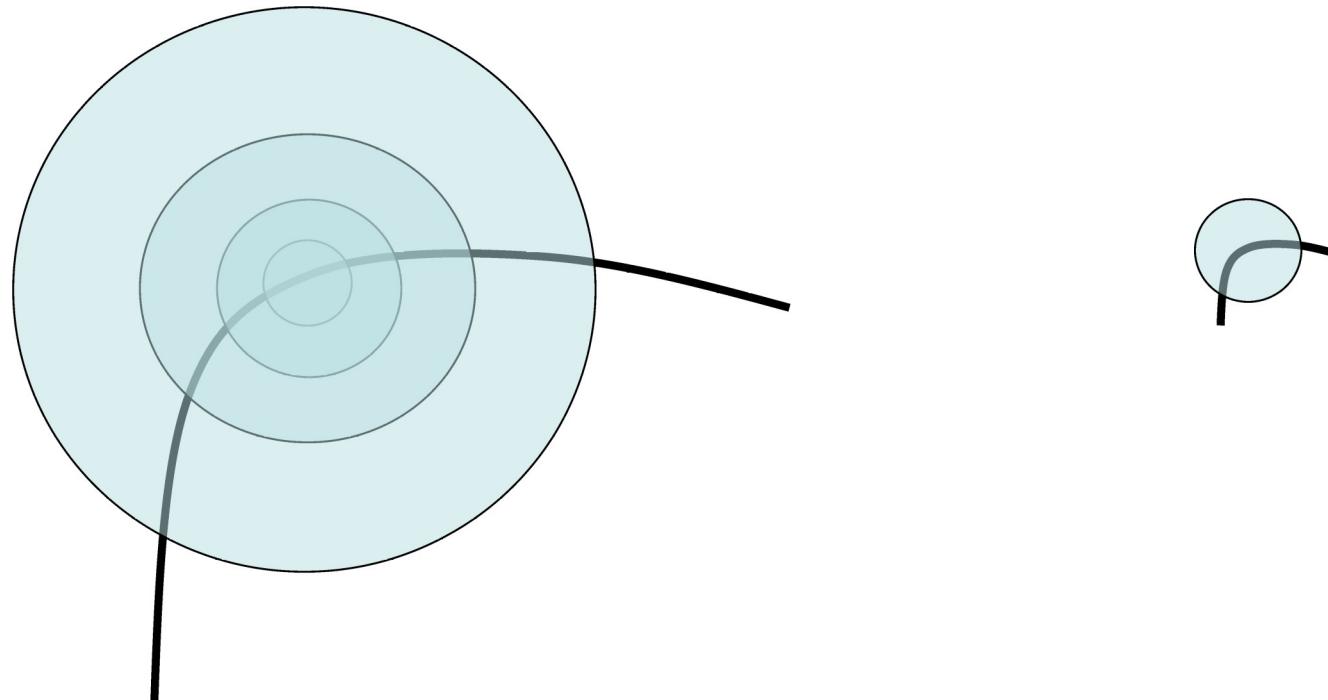
- The keypoints are highly dependent on window size



Scale invariant detector



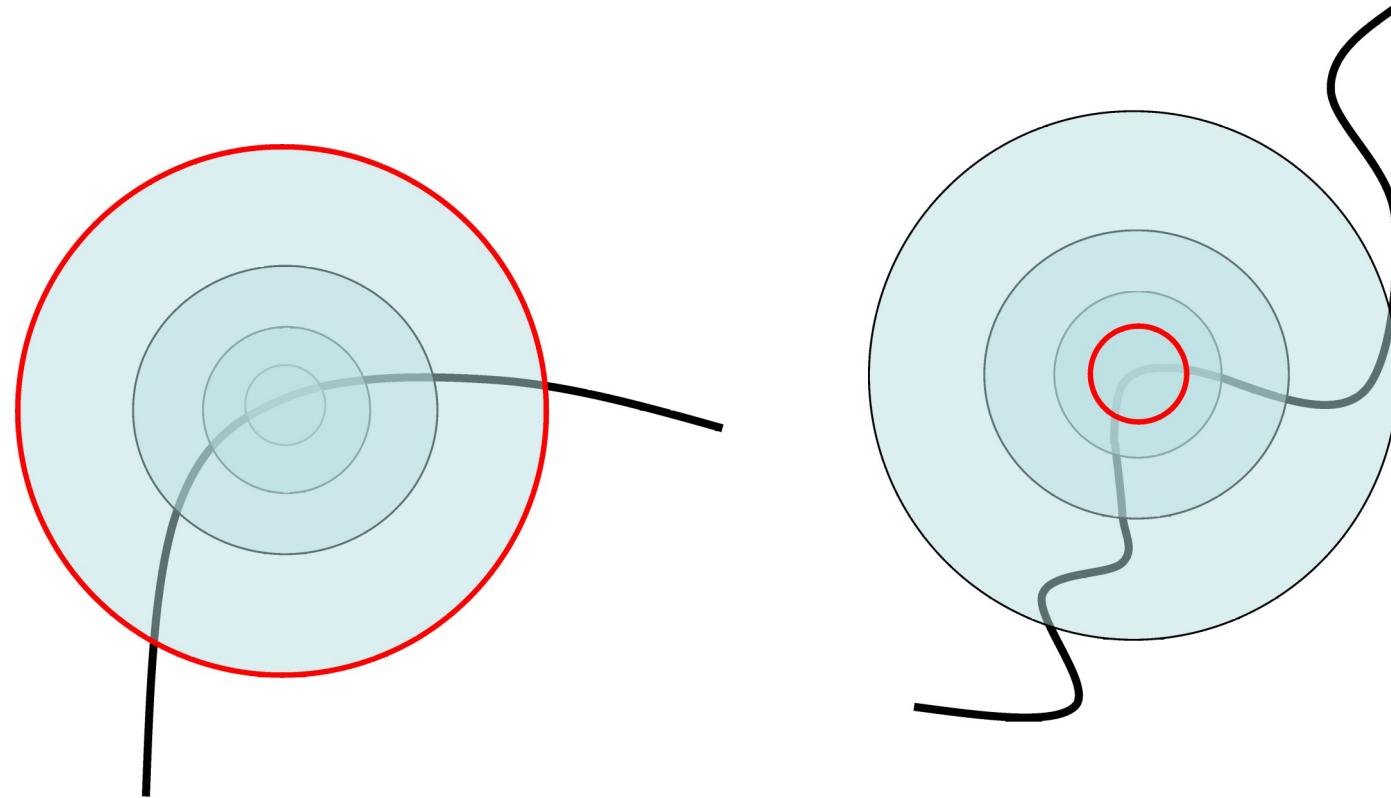
- Consider circles of different size around a point
- We can find a size that look similar in both images



Scale invariant detector



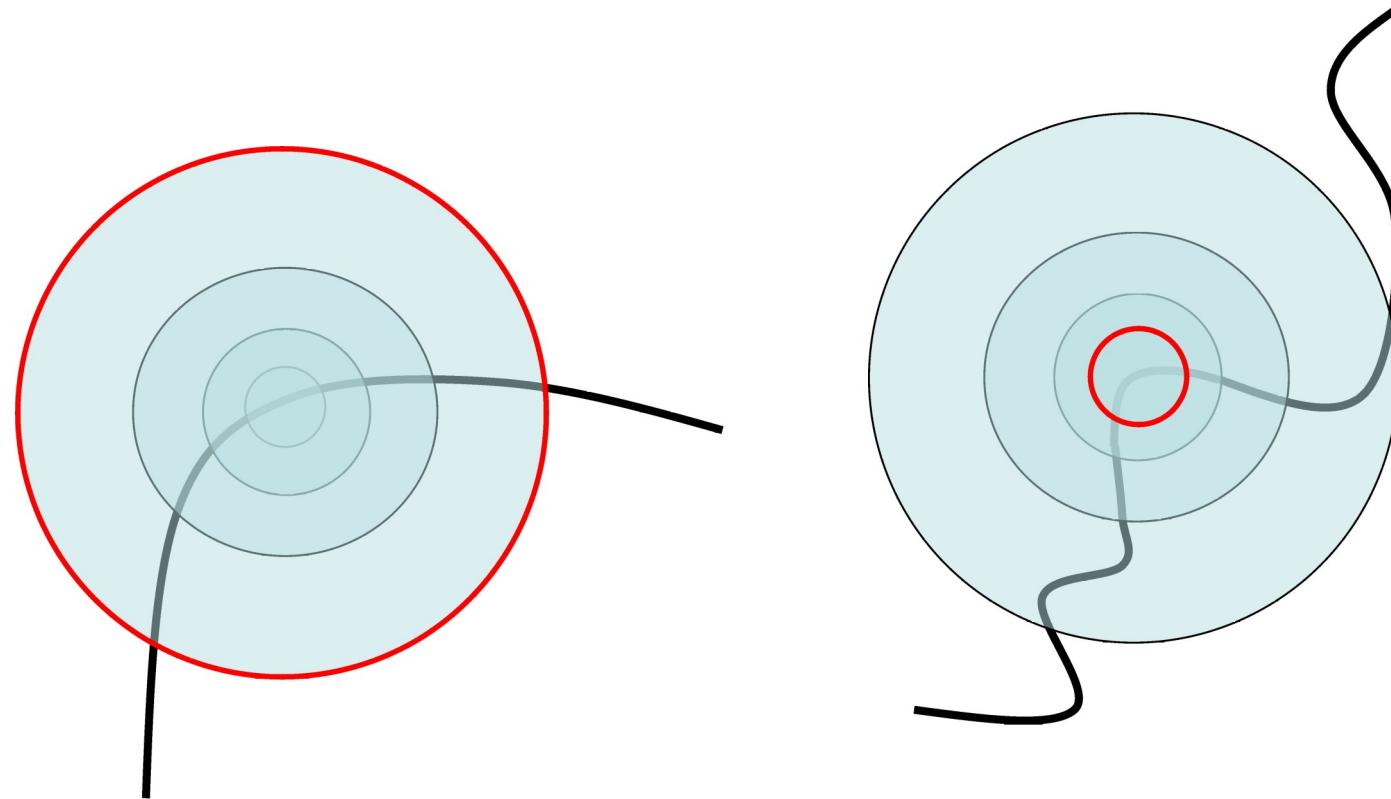
- How do we select appropriate circles size in both images?



Scale invariant detector



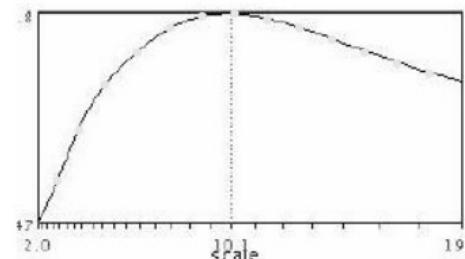
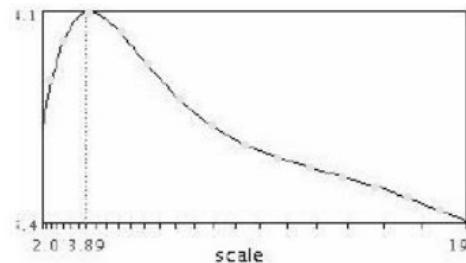
- We can select the window size that **maximize** Harris operator



Scale invariant detector



- Example



Scale invariant detector



- Instead computing θ for larger and larger windows a pyramidal approach can be used



(sometimes need to create in-between levels, e.g. a $\frac{3}{4}$ -size image)

Other scale invariant detectors

- **HARRIS-LAPLACIAN**
 - K. Mikolajczyk, C. Schmid, Indexing based on scale invariant interest points, ICCV 2001
- **SIFT:** Scale Invariant Feature Transform
 - D. Lowe. "Distinctive Image Features from Scale-Invariant Keypoints", IJCV 2004
- **FAST:** Features from Accelerated Segment Test
 - Edward Rosten and Tom Drummond, "Machine learning for high speed corner detection" in 9th European Conference on Computer Vision, vol. 1, 2006
- **SURF:** Speeded-Up Robust Features
 - Herbert Bay, Andreas Ess, Tinne Tuytelaars, Luc Van Gool, "SURF: Speeded Up Robust Features", Computer Vision and Image Understanding (CVIU), 2008
- **MSER:** Maximally Stable Extremal Regions
 - J. Matas, O. Chum, M. Urban, and T. Pajdla. "Robust wide baseline stereo from maximally stable extremal regions." Proc. of British Machine Vision Conference, 2002.
- **BRISK:** Binary Robust Invariant Scalable Keypoint
 - Stefan Leutenegger ; Margarita Chli ; Roland Y. Siegwart, 2011 International Conference on Computer Vision



UNIVERSITÀ DI PARMA

Features Extraction

Question time!

